

LLM Sensors — Server Deployment Guide

As of 29 September 2026 · Image `justfarbod/llm-sensors:research` · LLM backend: Ollama Cloud

This guide deploys the research platform on any Linux server or VM with Docker, Caddy for HTTPS, and Ollama Cloud for the models. Only creating the server, its firewall and its DNS name depend on the provider.

How to use this guide

- Paste every grey block into the terminal exactly as shown. Paste one block at a time and wait for it to finish.
- **SERVER** blocks run on the server over SSH. **YOUR COMPUTER** blocks run on your own machine.
- Only lines ending in `# <- edit` need your own value. Change the value after `=`, then paste.
- Blocks with a single `read` command wait for you to type or paste a value and press Enter. Always paste them on their own.
- If the terminal asks whether to paste several lines at once, confirm it.
- Every block is safe to run again if you are unsure whether it worked.

1. Overview

2. Requirements

3. Prepare the server

4. Ollama Cloud


5. Configuration files

6. Start and first-run setup

7. Browser telemetry extension

1. Overview

The platform runs as two Docker containers on any Linux server. Models come from Ollama Cloud over HTTPS, so the server needs no GPU and runs no local model.



Caddy is the only container reachable from the internet. It gets and renews a Let's Encrypt certificate on its own. The app container stores everything (SQLite database, uploads, experiment data) in the `llm-sensors-data` volume. That volume is the one thing you must back up.

Component	Image	Role
App	<code>justfarbod/llm-sensors:research</code>	Web UI, API, experiments, telemetry
Reverse proxy	<code>caddy:2</code>	HTTPS termination, certificates, websockets
LLM	Ollama Cloud (external)	Model inference via API key

2. Requirements

A 2 vCPU / 4 GB Linux server with a public IP, a DNS name and an Ollama account is enough. Inference happens at Ollama Cloud, so server size only depends on the app itself and the number of participants.

Item	Minimum	Recommended	Notes
App image	Built from commit <code>e5ed578f2</code> or later	Latest	Older images lack the extension name check and the Download extension button used in section 7

OS	64-bit Linux	Ubuntu Server 24.04 LTS	The commands in this guide assume Ubuntu or Debian
CPU architecture	Must match the image	x86-64 (amd64)	Images built on an x86 PC without <code>buildx --platform</code> are amd64 only
vCPU	2	2–4	
RAM	4 GB	4–8 GB	The app uses about 1–2 GB at rest
Disk	30 GB	64 GB SSD	The image is several GB; data grows with sessions and backups
Public IPv4	Static	Static	DNS and certificates break if it changes
DNS name	Any A record	e.g. <code>research.example.edu</code>	Needed for HTTPS; the extension is set to this exact hostname
Inbound ports	22, 80, 443	Same	80 is required for Let's Encrypt; do not open 8080
Outbound	HTTPS to the internet	Same	Docker Hub, Let's Encrypt, <code>ollama.com</code>
Software	Docker Engine + compose plugin	Same	Installed in section 3
LLM	Ollama account + API key	Same	See section 4

Ollama Cloud's free plan allows **1 concurrent request** and a limited set of starter models (ollama.com/pricing). This is fine for testing. For a study session with several participants at once, requests queue or fail. Pro (\$20/month) allows 3 and Max (\$100/month) allows 10.

Ollama states that prompts and responses are not logged or trained on. Confirm that sending participant chats to an external service is covered by your consent and data-handling approval.

3. Prepare the server

Install Docker, confirm the architecture, save your domain, and lock the firewall down to SSH and web traffic. Connect to the server over SSH as your normal sudo user.

3.1 Update the system

SERVER

```
sudo apt update && sudo DEBIAN_FRONTEND=noninteractive apt -y upgrade
```

3.2 Check the architecture

SERVER

It must print `x86_64` for an amd64 image.

```
uname -m
```

3.3 Install Docker

SERVER

Installs Docker Engine and the compose plugin with Docker's official script, then lets your user run Docker without `sudo`.

```
curl -fsSL https://get.docker.com | sh
sudo usermod -aG docker $USER
```

Log out with `exit` and connect with SSH again, so the group change applies. Then check that both commands print a version:

```
docker version && docker compose version
```

3.4 Enable the firewall

SERVER Skip this if your provider already has a network firewall that only allows 22, 80 and 443 (e.g. an Azure network security group).

```
sudo ufw allow OpenSSH
sudo ufw allow 80,443/tcp
sudo ufw --force enable
```

3.5 Save your domain

SERVER Point your DNS name at the server's public IP first. This block saves the name in the deployment folder, where every later step reads it, and prints two IP addresses: the one your name points to and the server's own. They must be the same.

```
mkdir -p ~/llm-sensors && cd ~/llm-sensors && touch .env && sed -i '/^DOMAIN=/d' .env
echo "DOMAIN=research.example.edu" >> .env # <- edit (hostname only, no https://)
. ./env && getent hosts "$DOMAIN"; curl -s https://ifconfig.me; echo
```

3.6 Turn on automatic security updates

SERVER Optional but recommended.

```
sudo DEBIAN_FRONTEND=noninteractive apt -y install unattended-upgrades
echo "unattended-upgrades unattended-upgrades/enable_auto_updates boolean true" \
| sudo debconf-set-selections
sudo dpkg-reconfigure -f noninteractive unattended-upgrades
```

4. Ollama Cloud

The app talks to <https://ollama.com> directly, authenticating with an API key. No Ollama install is needed on the server (docs.ollama.com/cloud).

1. Create an account at ollama.com and sign in.
2. Create an API key at ollama.com/settings/keys and copy it somewhere safe. It is shown only once. Don't confuse it with `WEBUI_SECRET_KEY` (step 5.1), which also looks like a long random string.

4.1 Test the key

SERVER Paste this line alone. It waits for you to paste the key and press Enter. The key is not shown on screen and not saved in your shell history.

```
read -s -p "Ollama API key: " OLLAMA_API_KEY; echo
```

SERVER Then send a test prompt. `gpt-oss:20b` is an example; use a model your plan includes.

```
curl -s -w "\nHTTP %{http_code}\n" https://ollama.com/api/chat \
-H "Authorization: Bearer $OLLAMA_API_KEY" \
-H "Content-Type: application/json" \
-d '{"model": "gpt-oss:20b", "messages": [{"role": "user", "content": "Hi"}], "stream": false}'
```

HTTP 200 with a `message.content` field means the key and model work. HTTP 401 means the key is wrong. An error about plan or credits means you need a model included in your plan.

Model names differ by client. Calling `ollama.com` directly (as the app does) uses names like `gpt-oss:20b`. The Ollama desktop app and CLI use a `-cloud` suffix such as `gpt-oss:20b-cloud`.

Monitor usage at ollama.com/settings/usage. The free plan's credits run out quickly with many participants.

5. Configuration files

Three files in `~/llm-sensors` define the whole deployment: `.env` for settings and secrets, `docker-compose.yml` for the containers, and `Caddyfile` for HTTPS. Each block below creates one file; nothing needs editing.

5.1 .env

SERVER Writes the settings next to the domain from step 3.5 and generates a random `WEBUI_SECRET_KEY`. Running it again keeps the existing domain, secret and extension settings.

```
cd ~/llm-sensors && . ./env
SECRET=$(grep -s '^WEBUI_SECRET_KEY=' .env | cut -d= -f2)
[ -n "$SECRET" ] || SECRET=$(openssl rand -hex 32)
EXT=$(grep '^EXPERIMENT_TELEMETRY_EXTENSION_' .env ||
echo EXPERIMENT_TELEMETRY_EXTENSION_ENABLED=false)
cat > .env <<EOF
DOMAIN=$DOMAIN
WEBUI_SECRET_KEY=$SECRET
OLLAMA_BASE_URL=https://ollama.com
ENABLE_OPENAI_API=false
ENABLE_SIGNUP=false
ENABLE_INITIAL_ADMIN_SIGNUP=true
$EXT
EOF
chmod 600 .env
grep -v SECRET .env
```

- `WEBUI_SECRET_KEY` signs login sessions. It must stay the same across red deploys. It is **not** the Ollama key.
- `ENABLE_OPENAI_API=false` hides the unused OpenAI connection, which would otherwise show errors.
- `ENABLE_SIGNUP=false` with `ENABLE_INITIAL_ADMIN_SIGNUP=true` lets exactly one person register, the admin. Participants are then created by the admin with bulk user generation.
- Settings such as `ENABLE_SIGNUP` and `OLLAMA_BASE_URL` are copied into the database on first start. Afterwards, change them in Admin Settings, because editing `.env` has no effect on them.

5.2 docker-compose.yml

SERVER The file is written in JSON form, which Docker Compose reads like YAML. Unlike YAML, it doesn't depend on indentation, so it still works if copying from the PDF drops leading spaces.

```
cd ~/llm-sensors
cat > docker-compose.yml <<'EOF'
{
  "services": {
    "llm-sensors": {
      "image": "justfarbod/llm-sensors:research",
      "container_name": "llm-sensors",
      "restart": "unless-stopped",
      "env_file": ".env",
      "environment": {
        "WEBUI_URL": "https://${DOMAIN}",
        "CORS_ALLOW_ORIGIN": "https://${DOMAIN}"
      },
      "volumes": ["llm-sensors-data:/app/backend/data"],
      "logging": {"driver": "json-file", "options": {"max-size": "20m", "max-file": "5"}}
    },
    "caddy": {
      "image": "caddy:2",
      "container_name": "caddy",
      "restart": "unless-stopped",
      "ports": ["80:80", "443:443"],
      "environment": {"DOMAIN": "${DOMAIN}"},
      "volumes": [
        "./Caddyfile:/etc/caddy/Caddyfile:ro",
        "caddy-data:/data",
        "caddy-config:/config"
      ],
      "depends_on": ["llm-sensors"]
    }
  },
  "volumes": {"llm-sensors-data": {}, "caddy-data": {}, "caddy-config": {}}
}
EOF
docker compose config --quiet && echo "docker-compose.yml is valid"
```

The app container publishes no ports, so it is only reachable through Caddy. This matters because Docker-published ports bypass `ufw`. The `research` tag pins the image you built.

5.3 Caddyfile

SERVER Caddy obtains the certificate for `DOMAIN` automatically and proxies websockets, which chat streaming needs, without extra settings.

```
cd ~/llm-sensors
cat > Caddyfile <<'EOF'
${DOMAIN} {
  encode zstd gzip
  reverse_proxy llm-sensors:8080
}
EOF
```

6. Start and first-run setup

Pull and start the stack, register the admin account before anyone else can, then add the Ollama Cloud key. The whole process takes about 10 minutes.

6.1 Start the containers

SERVER The last line prints your site's address.

```
cd ~/llm-sensors
docker compose pull
docker compose up -d
. ./env && echo "Open https://$DOMAIN in 2-3 minutes"
```

The first start runs database migrations and requests the HTTPS certificate, which takes a few minutes. Until then the browser may show an error; reload after a minute.

6.2 Set up the platform in the browser

1. Open the printed address and **register the admin account immediately**. The first account becomes admin. With this `.env`, sign-up closes after it.
2. Add the Ollama Cloud key: open **Admin Settings** → **Connections**, and under **Ollama API** click the settings icon next to `https://ollama.com`.
 - Set **Auth** to **Bearer** and paste the Ollama key into **API Key**.
 - Optional: under **Model IDs**, list only the models participants may use (e.g. `gpt-oss:20b`). Otherwise every cloud model appears in the model menu.
 - Click **Verify Connection**, then **Save** in the dialog and **Save** again on the Connections page.
3. Open a new chat, pick the model from the menu at the top and send a message. A streamed reply confirms the whole path (browser, Caddy, websockets, app, Ollama Cloud).
4. Confirm sign-ups are off: **Admin Settings** → **General** → **Enable New Sign Ups** should be off.
5. Create participants: **Users** → bulk user generation, then export the generated credentials. Put them into groups under **Users** → **Groups**.
6. Set up the study: import or build a workflow under **Experiment Tasks** and apply it to a group. See `ADMIN_PANEL_GUIDE.md` sections 4–7 in the repository.

7. Browser telemetry extension

Participants receive the extension as a zip file from a download link and load it into Chrome themselves. You start from the ready-made base zip, set your server's hostname in it, and zip it again. No build or repository is needed. The server recognizes the extension by its name, so there is no ID to copy. Skip this section if the study does not use the extension; `.env` already keeps it disabled.

The extension records tab URLs and titles, keystrokes, and copied, cut and pasted text in tracked fields. Consent materials must say so (see `browser-extension/experiment-telemetry/README.md` in the repository).

The base zip is at drive.google.com/file/d/1RxEBmyuhs0WZdvRaFn0r4bAQiDIQUYy. It contains version 2.2.0, set to `llm-sensors.belgiumcentral.cloudapp.azure.com`. If that is your server's hostname, skip step 7.1 and use the base zip as it is.

7.1 Set your hostname in the zip

Use either the command (option A) or edit the file by hand (option B). Both produce `llm-sensors-extension.zip` for your server.

Option A: one command

YOUR COMPUTER Run in a bash terminal (Linux, macOS or WSL; needs `curl` and `python3`). Downloads the base zip, sets the hostname and writes the new zip to `~/llm-sensors-extension`.

```
DOMAIN=research.example.edu # <- edit (same hostname as on the server)
mkdir -p ~/llm-sensors-extension && cd ~/llm-sensors-extension
curl -fsSL -o base.zip \
"https://drive.google.com/uc?export=download&id=1RxEBmyuhs0WZdvRaFn0r4bAQiDI1QUYy"
python3 - "$DOMAIN" <<'EOF'
import json, sys, zipfile
origin = "https://" + sys.argv[1].replace("https://", "").strip("/")
src = zipfile.ZipFile("base.zip")
files = {name: src.read(name) for name in src.namelist()}
manifest = json.loads(files["manifest.json"])
manifest["host_permissions"] = [origin + "/*"]
for script in manifest["content_scripts"]: script["matches"] = [origin + "/*"]
manifest["action"]["default_title"] = "Connect to your experiment at " + origin
files["manifest.json"] = json.dumps(manifest, indent=2)
out = zipfile.ZipFile("llm-sensors-extension.zip", "w", zipfile.ZIP_DEFLATED)
[out.writestr(name, data) for name, data in files.items()]
out.close()
print("Created ~/llm-sensors-extension/llm-sensors-extension.zip for " + origin)
EOF
```

YOUR COMPUTER Optional, on WSL only: opens the folder in Windows so you can find the zip.

```
explorer.exe "$(wslpath -w ~/llm-sensors-extension)"
```

Option B: by hand

1. Download the base zip from the link above and unzip it into a new folder.
2. Open `manifest.json` from that folder in a text editor (e.g. Notepad or TextEdit).
3. Use **Find and Replace** to replace `llm-sensors.belgiumcentral.cloudapp.azure.com` with your hostname. It appears **3 times**: under `host_permissions`, under `matches`, and in `default_title`. Keep the `https://` in front and the `/*` after it. Save the file.
4. Open the folder, select **all the files inside it** (not the folder itself) and compress them into a zip: on Windows, right-click → **Compress to ZIP file** (or **Send to** → **Compressed (zipped) folder**); on macOS, right-click → **Compress**. Rename the result to `llm-sensors-extension.zip`.

To move to another hostname later, repeat step 7.1 with the new name.

7.2 Upload the zip

Upload `llm-sensors-extension.zip` to any file host that gives participants an HTTPS link, e.g. Google Drive with sharing set to **Anyone with the link**, or your university cloud storage. This is the zip participants download; they don't change anything in it.

7.3 Enable the extension on the server

SERVER Paste this line on its own, then paste the download link from step 7.2 and press Enter:

```
read -p "Download link (https://...): " EXT_URL
```

SERVER Then write the extension settings and apply them. Data stays in the volume.

```
cd ~/llm-sensors && . ./env
sed -i '/^EXPERIMENT_TELEMETRY_EXTENSION_/d' .env
cat >> .env <<EOF
EXPERIMENT_TELEMETRY_EXTENSION_ENABLED=true
EXPERIMENT_TELEMETRY_EXTENSION_ORIGIN=https://$DOMAIN
EXPERIMENT_TELEMETRY_EXTENSION_NAME="Open WebUI Experiment Telemetry"
EXPERIMENT_TELEMETRY_EXTENSION_STORE_URL="$EXT_URL"
EXPERIMENT_TELEMETRY_EXTENSION_MIN_VERSION=2.2.0
EOF
docker compose up -d
```

The experiment Start screen now shows a **Download extension** button that opens your download link. If you later change the download link or hostname, repeat this step with the new values.

7.4 Installation steps for participants

Send participants these steps, and follow them once yourself in a participant account before the study.

1. On the experiment Start screen, click **Download extension** and download the zip. On a Google Drive page, use its **Download** button.
2. Unzip the file into a folder you keep, e.g. in Documents. Chrome loads the extension from this folder, so don't delete or move it during the study.
3. In Chrome, open `chrome://extensions`, turn on **Developer mode** (top right), click **Load unpacked** and select the unzipped folder (the one containing `manifest.json`).
4. Return to the experiment tab and wait for **Extension connected**. There is no need to click the extension icon.